

---

# Transight Hackathon Problem Statement

## Multimodal Omni-Channel Conversation Intelligence Backend

---

### Background

Modern enterprises—especially in **banking, telecom, and customer support domains**—handle large volumes of customer interactions across **voice calls and digital text channels**, often in **multiple languages**. Deriving actionable insights from these conversations is critical for improving service quality, ensuring compliance, and managing operational risk.

With the availability of **multimodal AI models capable of directly understanding audio and text**, it is now possible to build intelligent backend systems that can analyze conversations without traditional speech-to-text or language-specific pipelines.

This hackathon challenges participants to design and implement a **backend-only, API-driven conversation intelligence system** that can process **voice or text conversations** and generate **structured, configurable insights** suitable for real-world enterprise integration.

---

### Problem Statement

**Design and implement a backend API system that can analyze multimodal customer conversations (voice or text) across languages and generate structured insights based on configurable client context and business rules.**

The solution must be API-first and suitable for integration into mobile or web applications.

---

### Scope & Functional Requirements

#### 1. Input Handling (Multimodal)

- The system must accept **either**:
  - An audio file (customer support call recording), **or**
  - A text-based conversation transcript
- Language should be **auto-detected** by the system.

---

## 2. Core Conversation Intelligence (Mandatory)

Using a multimodal AI model, extract the following insights:

- Conversation summary
  - Detected language(s)
  - Overall sentiment or emotional tone
  - Primary customer intent(s)
  - Key topics or entities discussed
- 

## 3. Configurable Client Context (Mandatory)

The system must support **client-defined configuration** that influences analysis.

Example (JSON-based configuration):

- Business domain (e.g., banking, telecom)
- Products or services
- Policies or rules
- Risk or compliance triggers

Participants may implement this as:

- JSON input
  - Static configuration file
  - Simple database schema
- 

## 4. Advanced Analysis (Choose at Least One)

Implement **at least one** of the following:

- Compliance or policy violation detection
  - Agent quality or performance scoring
  - Call outcome classification (resolved, escalated, dropped, etc.)
  - Risk or escalation score generation
- 

## 5. Output Requirements

- Output must be a **well-structured JSON response**
- Designed as if it will be consumed by an enterprise application
- Must include both:

- Analytical results
  - Any detected risks, flags, or classifications
- 

## 6. Technical Expectations

- Backend-only solution (no UI required)
  - Clear API endpoints
  - Well-documented request/response formats
  - Clean, readable, and maintainable code
  - README explaining setup, design decisions, and assumptions
- 

## Optional Enhancements (Bonus)

- Multiple speakers handling
  - Timeline-based sentiment or emotion tracking
  - Extensible schema for adding new insight types
  - Basic authentication or API key handling
- 

## Deliverables

1. Source code repository
  2. API documentation (OpenAPI / Markdown)
  3. Sample requests and responses
  4. README explaining:
    - Architecture
    - AI usage approach
    - Configuration mechanism
    - Limitations and future improvements
- 

## Shortlisting & Evaluation Rubric

Category	Weight	What We Evaluate
<b>Problem Understanding</b>	20%	Clarity of interpretation, alignment with requirements, logical approach
<b>API Design &amp; Backend Architecture</b>	20%	Endpoint design, request/response schemas, modularity, scalability thinking

<b>AI Reasoning &amp; Insight Quality</b>	25%	Relevance, accuracy, and usefulness of extracted insights
<b>Configurability &amp; Flexibility</b>	15%	How cleanly client context or rules influence analysis
<b>Code Quality &amp; Documentation</b>	10%	Readability, structure, comments, and README clarity
<b>Innovation &amp; Practicality</b>	10%	Real-world usefulness, thoughtful extensions, enterprise readiness

---

## Intern Selection Criteria

Candidates demonstrating the following will be prioritized for the **AI Internship Program**:

- Strong problem decomposition and reasoning skills
- Ability to design production-ready backend APIs
- Practical application of multimodal AI models
- Clear thinking around enterprise use cases
- Clean code and documentation practices

## Final Note to Participants

This challenge is not about building a perfect product in limited time—it is about **how you think, design, and apply AI to solve real enterprise problems**. Simplicity, clarity, and correctness are valued over complexity.

*Wish you All the very best* 